



**DATA SCIENCE
LAB**



INFORMATION
TECHNOLOGY
UNIVERSITY

Data Engineering



What is Data Engineering?

“Data engineering comprises all engineering and operational tasks required to make data available for the end-user, whether for the purpose of analytics, model building or app development etc.”

3 step process in layman's terms.

1. Taking raw data.
2. Doing a bunch of work to it.
3. Delivering a clean dataset of database.



What do Data Engineers do?

- Model data
- Build production ready data warehouses and data lakes.
- Tools that process massive amount of data.
- Automate data pipelines.

THE DATA SCIENCE HIERARCHY OF NEEDS


LEARN/OPTIMIZE

AGGREGATE/LABEL

EXPLORE/TRANSFORM

MOVE/STORE

COLLECT



THE DATA SCIENCE HIERARCHY OF NEEDS

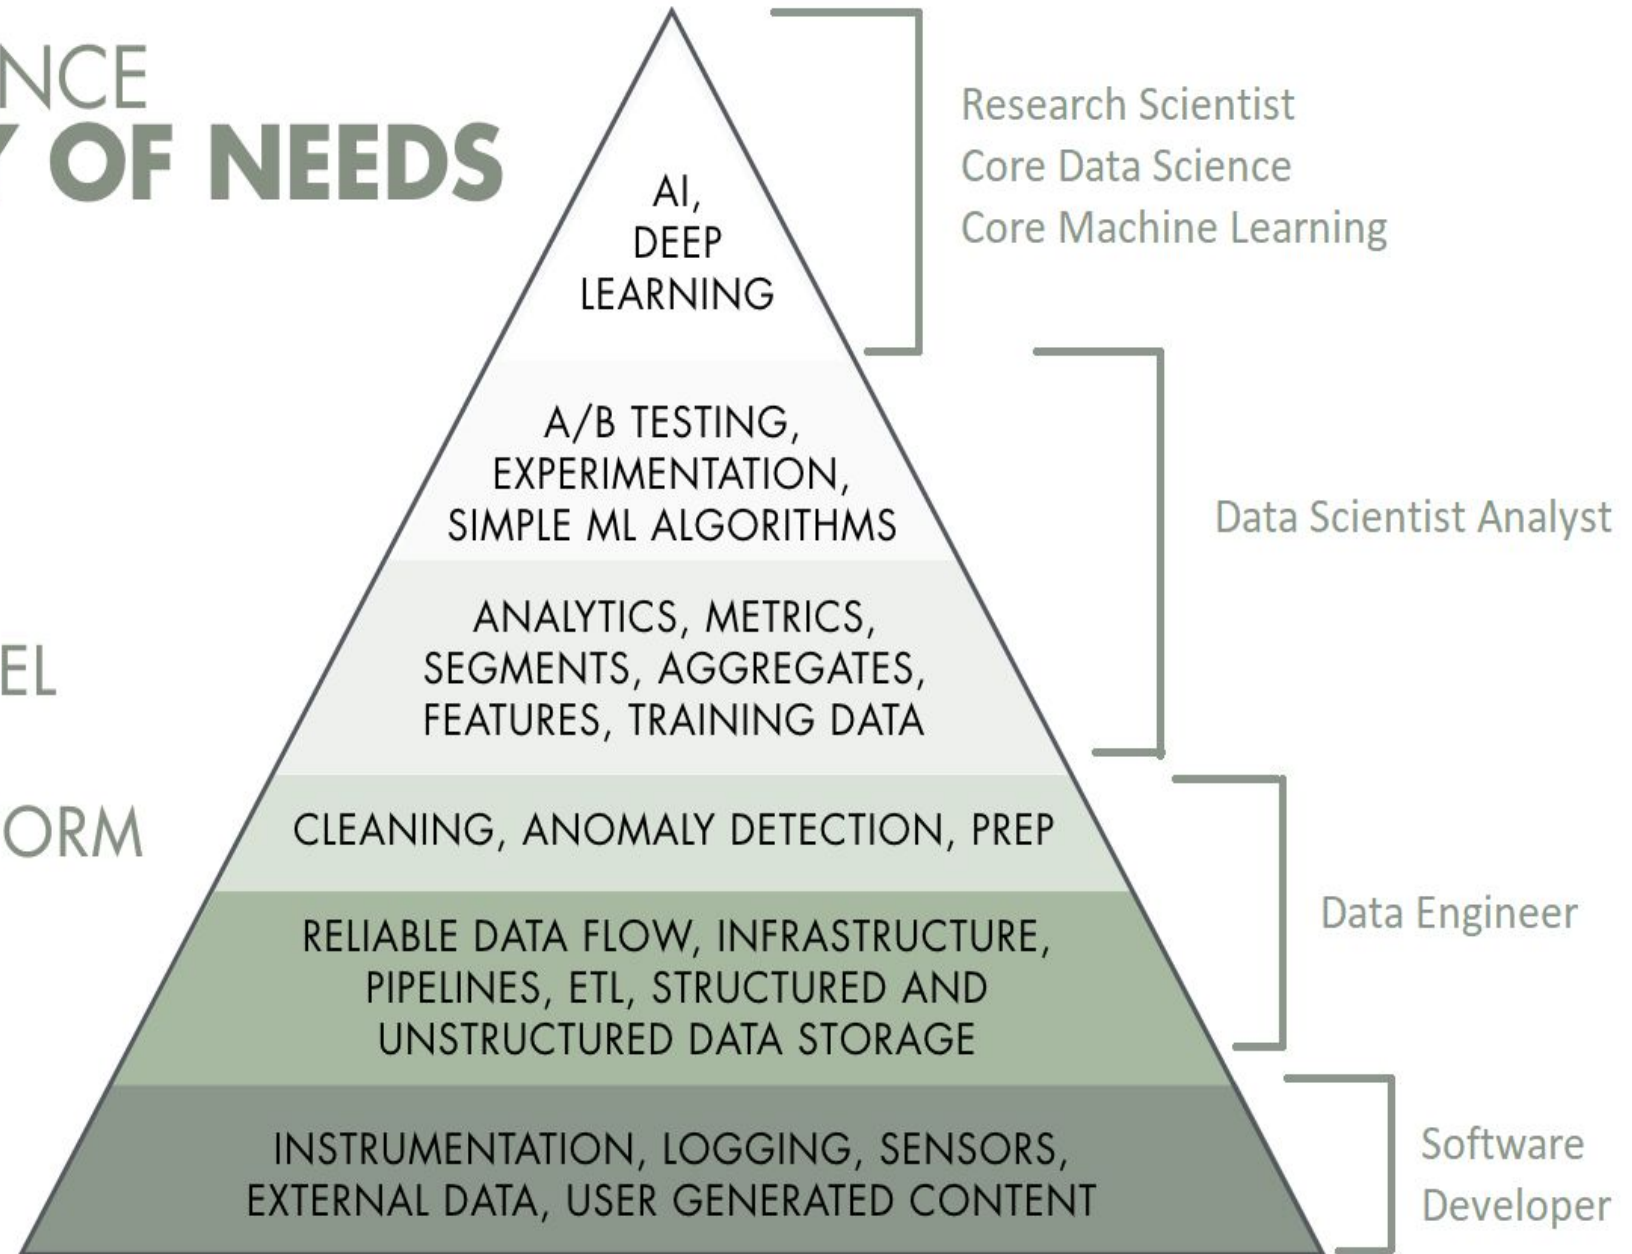
LEARN/OPTIMIZE

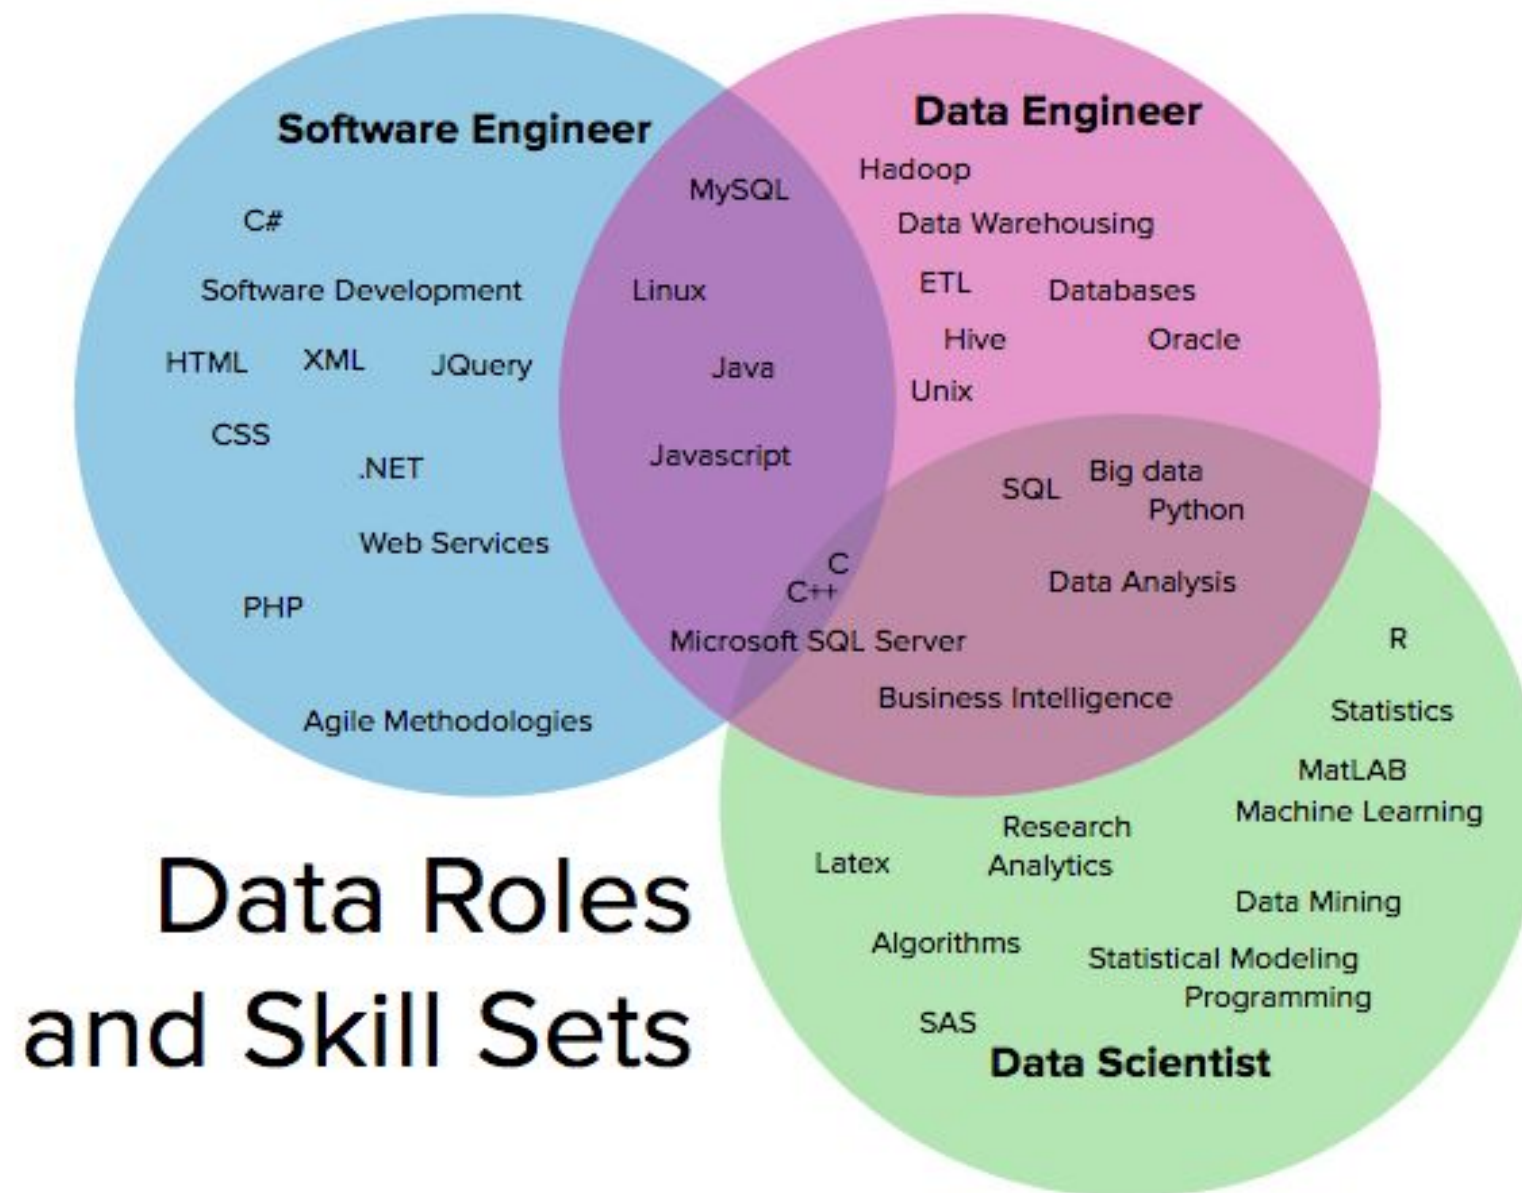
AGGREGATE/LABEL

EXPLORE/TRANSFORM

MOVE/STORE

COLLECT





Common data engineering activities.

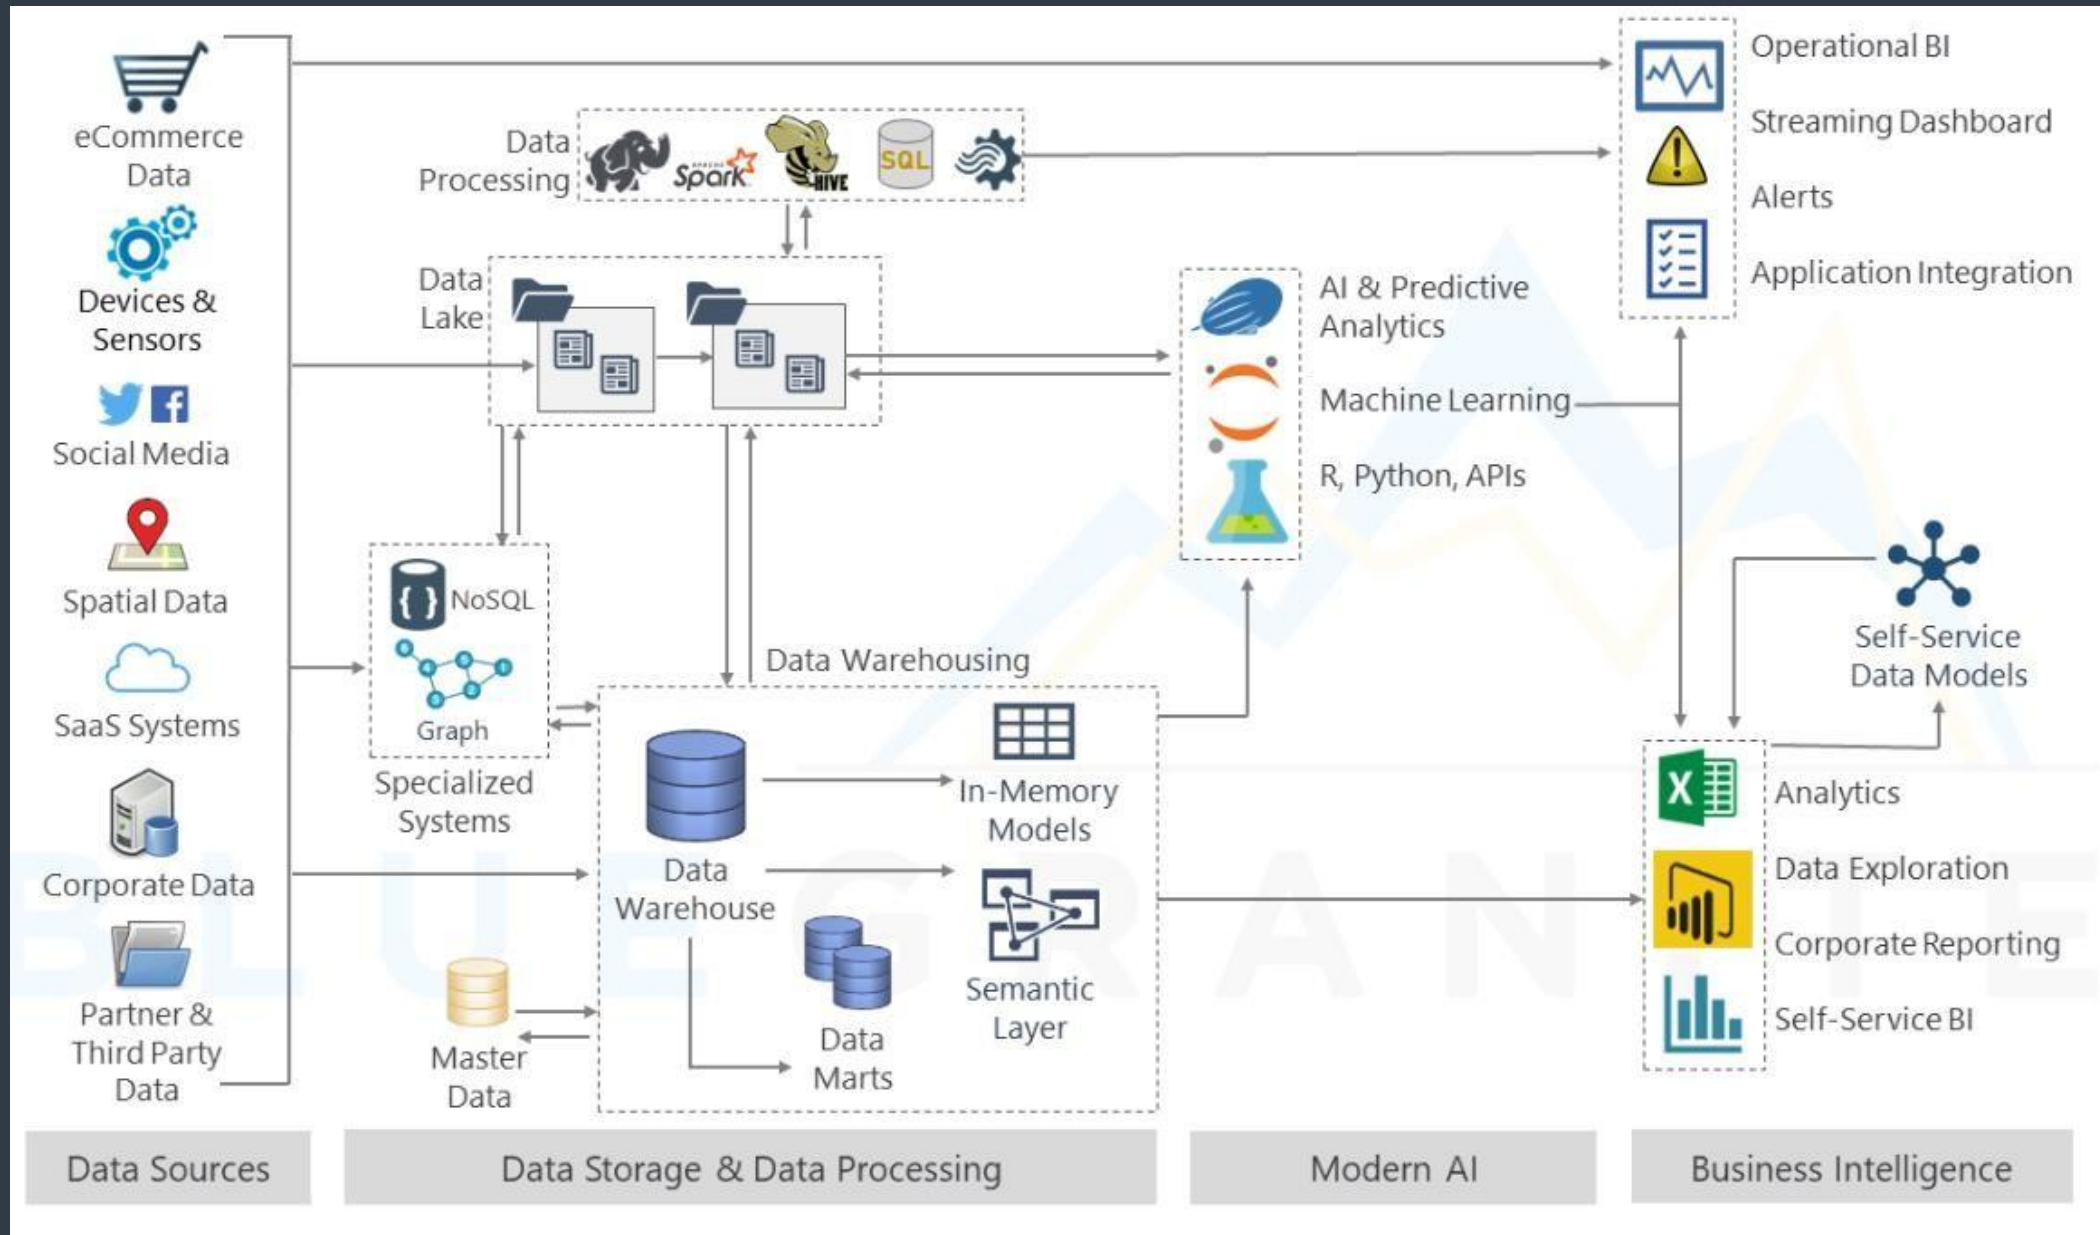


1. Ingest data from a data source.
2. Build and maintain data warehouse.
3. Create a data pipeline.
4. Create an analytics table for a specific use case.
5. Migrate data to the cloud.
6. Schedule and automate pipelines.
7. Debug data quality issues.
8. Optimize Queries.
9. Design a database.



Data Modeling

Data Architecture



Data Modeling

“Data Modeling is an abstraction that organizes elements of data and how they will relate to each other”

– Wikipedia

Example: Spreadsheets for household

- You define rows and columns
- You structure your data


Process of Data Modeling

The process of data modeling is to

- Organize data into databases.
- To ensure that your data is persistent.
- To ensure that it is easily useable by you and your organization.

Data Modeling is also called **database modeling**.

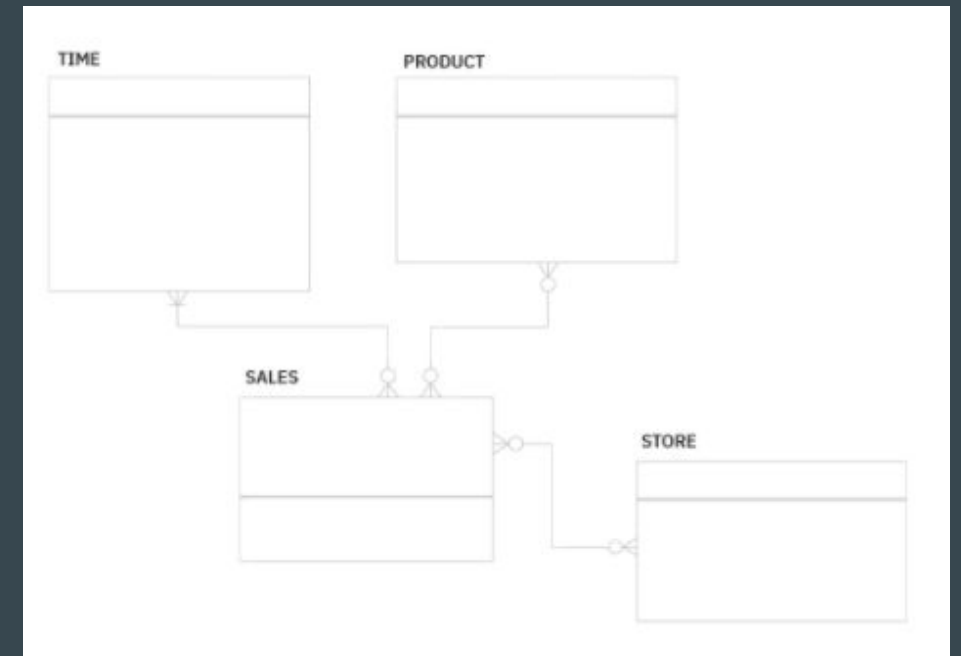
Data Modeling



- Process to support business and user applications
- Gather requirements
- Conceptual Data Modeling
- Logical Data Modeling
- Physical Data Modeling

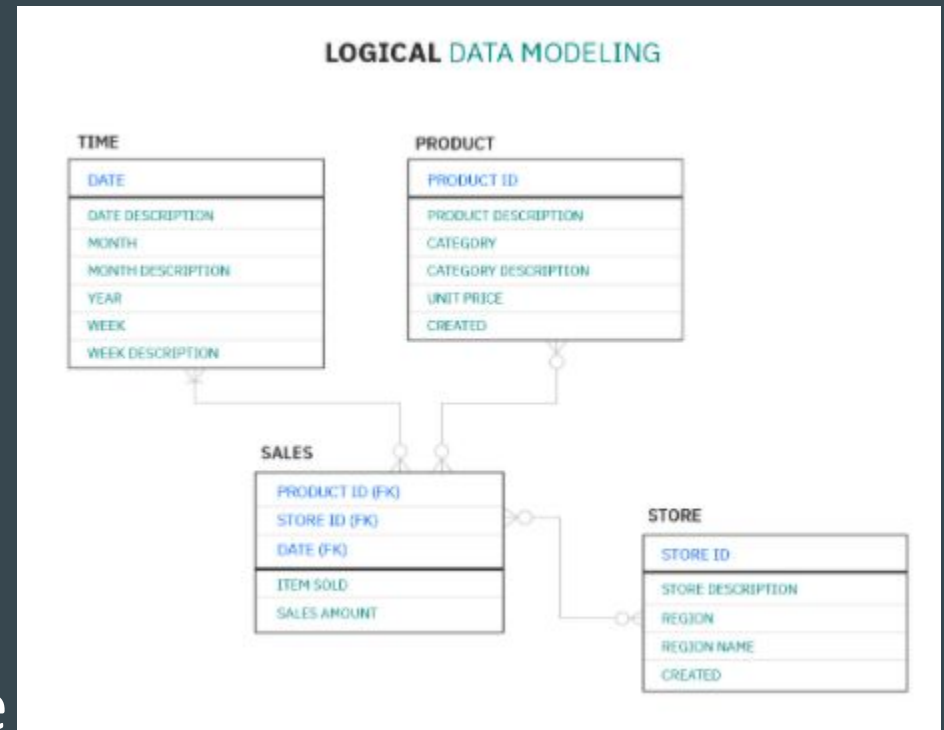
Conceptual Data Modeling

- Offers a big view picture of the business structure
- Created as part of the process of gathering initial project requirements
- Typically includes **entity classes**, their **characteristics** and **constraints** and the **relationships** between them



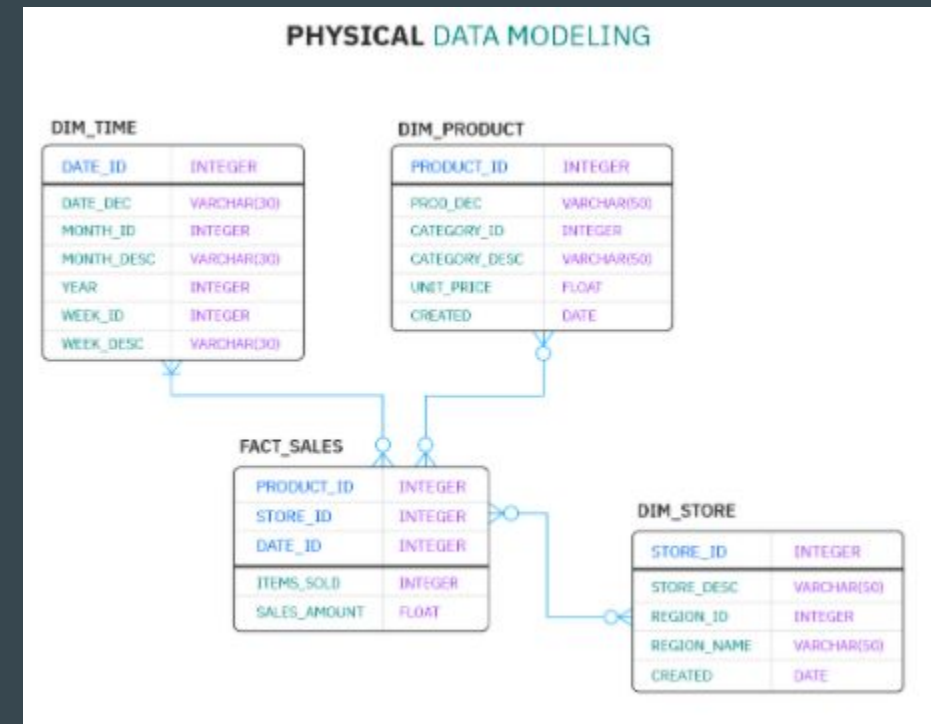
Logical Data Modeling

- Greater detail about the system
- More concerned about system implementation
- Data attributes in each entity are defined
- Data attributes, such as **data types** and **lengths** and relationships between **entities** are indicated



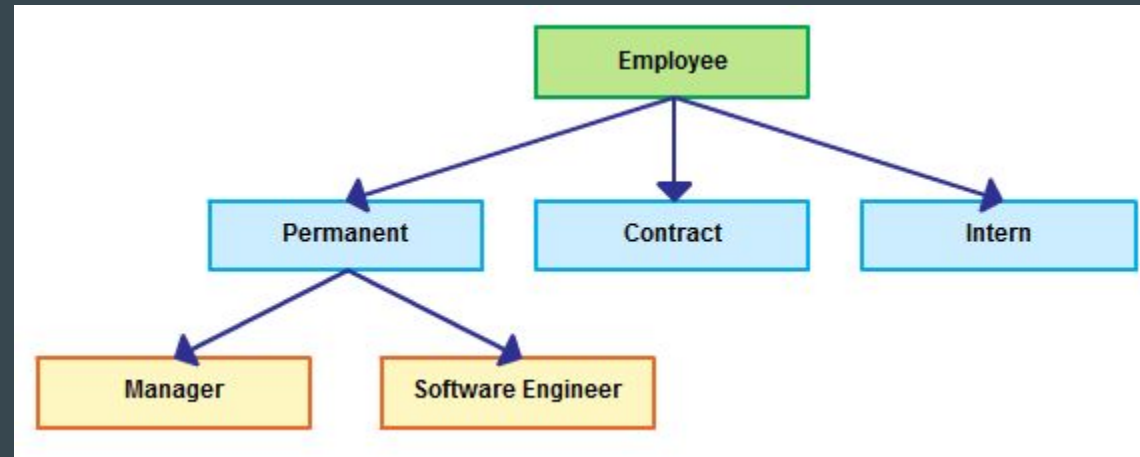
Physical Data Modeling

- Demonstrates the low-level implementation details
- A finalized design is offered containing data types, primary and foreign keys
- Can include DBMS-specific properties, including performance tuning.



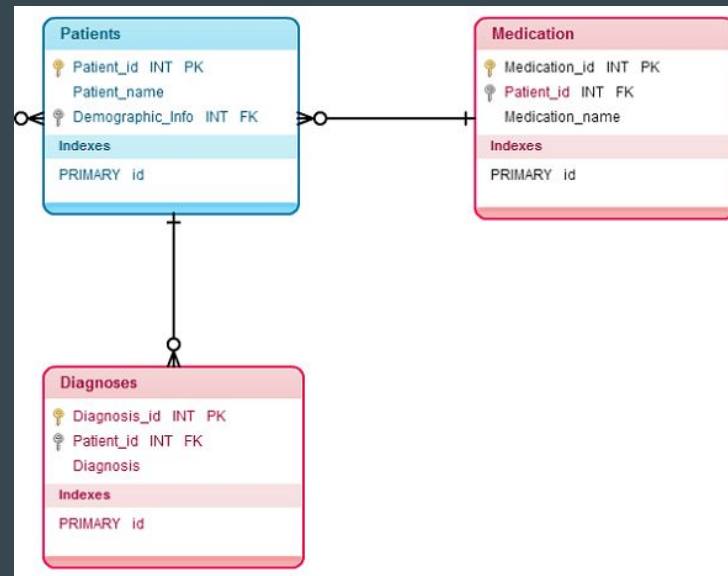
Types of Data Modeling

- Hierarchical Data Models
 - Relationships represented in a tree-like format
 - Each record has a single root/parent and maps to child tables



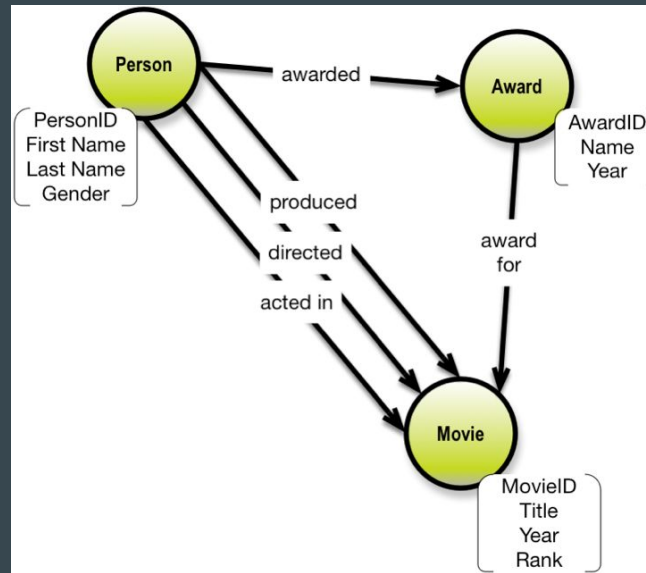
Types of Data Modeling

- Relational Data Models
 - Data segments are explicitly joined through the use of tables, reducing database complexity.



Types of Data Modeling

- Graph Data Models
 - Based on Graph Theory
 - Nodes and Edges in a graph are used to represent data



Why is data modeling important?




- Data organization is critical
- Organized data determines later data use
- Begin prior to building out application, business logic, and analytical models
- Iterative process



Relational and NoSQL Databases

Relational Model

"This model **organizes data into** one or more tables (or "**relations**") **of columns and rows**, with a **unique key identifying each row**. Generally, each table represents one "entity type" (such as customer or product)."



Relational Database

Invented by Edgar Codd (1970)

"... is a digital database **based on the relational model** of data...a software system used to maintain relational databases is a relational database management system (RDBMS)."

-- Wikipedia

Relational Database

"SQL (Structured Query Language) is the language used across almost all relational database system for querying and maintaining the database."

-- Wikipedia

Common Types of Relational Databases

- Oracle
- Teradata
- MySql
- PostgreSQL
- Sqlite

The Basics

- Database/Schema
 - Collection of Tables
- Tables/Relation
 - A group of rows sharing the same labeled elements
 - Customers

Employee

Name	Empld	DeptName
Harry	3415	Finance
Sally	2241	Sales
George	3401	Finance
Harriet	2202	Sales

Dept


DeptName	Manager
Finance	George
Sales	Harriet
Production	Charles

Customers

Name	Email	City
Amanda	jdoe@xyz.com	NYC
Toby	n/a	NYC

The Basics

- Columns/Attribute
 - Labeled element
 - Name, email, city
- Rows / Tuple
 - A single item
 - Amanda,
jdoe@xyc.com,
NYC



Customers

Name	email	City
Amanda	jdoe@xyz.com	NYC
Toby	n/a	NYC

Advantages of using a Relational Database

- Ease of use -- SQL
- Ability to do JOINS
- Ability to do aggregations and analytics
- Smaller data volumes
- Easier to change business requirements
- Flexibility for queries
- Modeling the data not modeling queries
- Secondary Indexes available
- ACID Transactions --data integrity
-

ACID Properties (Atomicity, Consistency, Isolation, Durability)

“...properties of database transactions intended to **guarantee validity** even **in the event of errors, power failures...**”

-- Wikipedia

Atomicity

“...the whole transaction is processed or nothing is processed”
-- Wikipedia

Consistency

"...only transactions that abide by constraints and rules is written into the database otherwise database keeps previous state"

-- Wikipedia

Isolation

"...transactions are processed independently and securely, order does not matter"

-- Wikipedia

Durability

“...completed transactions are saved to database even of cases of system failure”

-- Wikipedia



When Not to use a Relational Database?

When to not use a Relational Database

- Large amounts of data
- Need to be able to store different data type formats
- Need high throughput -- fast reads
- Need a flexible schema
- Need high availability
- Need horizontal scalability

PostgreSQL

What is PostgreSQL?

- Open source object-relational database system
- Uses and builds on SQL language



PostgreSQL

PostgreSQL Pros and Cons

Pros

- This database management engine is scalable and can handle terabytes of data.
- It supports JSON.
- There are a variety of predefined functions.
- A number of interfaces are available.

Cons

- Documentation can be spotty, so you may find yourself searching online in an effort to figure out how to do something.
- Configuration can be confusing.
- Speed may suffer during large bulk operations or read queries.


Reference: <https://www.keycdn.com/blog/popular-databases>

Comparison of Postgres with SQLite and MySQL

Name	SQLite	MySQL	PostgreSQL
Architecture	File Based	Client Server	Client Server
Transactional consistency	ACID	ACID	ACID
Replication	None	Master-Slave Replication, Master-Master Replication	Master-Slave Replication
Programming Language (Base Code)	C, C++	C, C++	C
Popular Use-Cases	Low-Medium Traffic Websites, IoT and Embedded Devices, Testing and Development	Web Sites, Web Applications, LAMP stack, OLTP-based applications	Analytics, Data Mining, Data Warehousing, Business Intelligence, Hadoop
Key Customers	Adobe, Facebook, and Apple	GitHub, Facebook, and YouTube	Cloudera, Instagram, and ViaSat

Reference: <https://logz.io/blog/relational-database-comparison/>

Famous Companies using PostgreSQL



UBER



Spotify®



reddit

Questions?