

Introduction to HBase

What is HBase?

HBase is a distributed
column-oriented data store
built on top of HDFS

What is HBase?

Apache HBase™ is the Hadoop database, a distributed, scalable, big data store




What is HBase?

Wikipedia:

HBase is an open source, non-relational, distributed database modeled after Google's BigTable and written in Java

Hadoop Ecosystem

The Hadoop Ecosystem



HBase is built on top of HDFS

HBase files are internally stored in HDFS

Why HBase when we have HDFS?

HBase vs. HDFS

HDFS is good for batch processing (scans over big files)

Not good for record lookup

Not good for incremental addition of small batches

Not good for updates

HBase vs. HDFS

HBase is designed to efficiently address

- Random Access

- Fast record lookup

- Support for record-level insertion


- Support for updates


HDFS	HBase
HDFS is a distributed file system great for storing large files	HBase is a database built on top of the HDFS
HDFS does not support fast individual record lookups	HBase provides fast lookups for large tables
High latency batch processing; no concept of batch processing	Low latency access to single rows from billions of records; enables Random access
Sequential access of data only	Internally uses Hash tables and provides random access; stores the data in indexed HDFS files for faster lookups

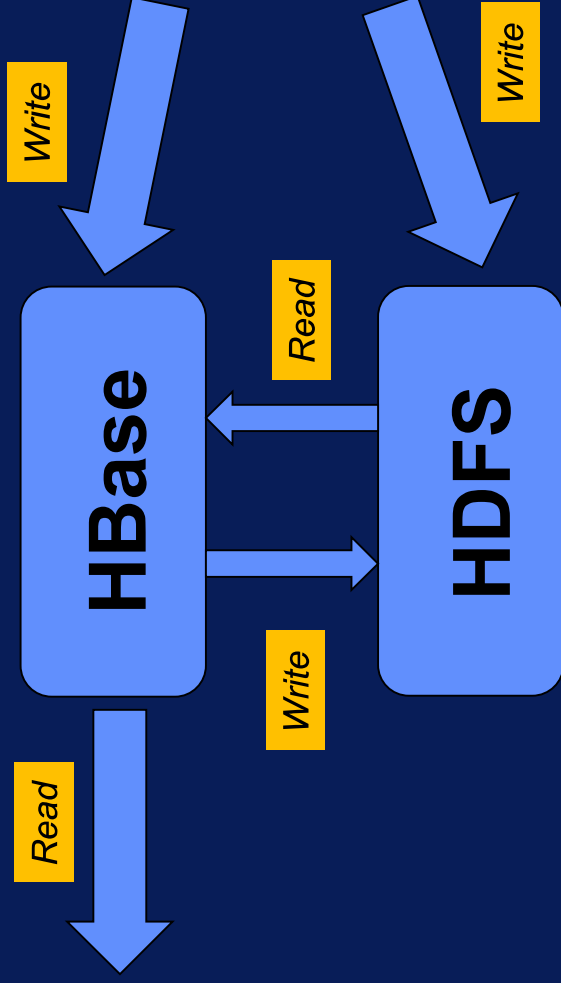
HDFS	HBase
HDFS is a distributed file system great for storing large files	HBase is a database built on top of the HDFS
HDFS does not support fast individual record lookups	HBase provides fast lookups for large tables
High latency batch processing; no concept of batch processing	Low latency access to single rows from billions of records; enables Random access
Sequential access of data only	Internally uses Hash tables and provides random access; stores the data in indexed HDFS files for faster lookups

HDFS	HBase
HDFS is a distributed file system great for storing large files	HBase is a database built on top of the HDFS
HDFS does not support fast individual record lookups	HBase provides fast lookups for large tables
High latency batch processing; no concept of batch processing	Low latency access to single rows from billions of records; enables Random access
Sequential access of data only	Internally uses Hash tables and provides random access; stores the data in indexed HDFS files for faster lookups

HDFS	HBase
HDFS is a distributed file system great for storing large files	HBase is a database built on top of the HDFS
HDFS does not support fast individual record lookups	HBase provides fast lookups for large tables
High latency batch processing; no concept of batch processing	Low latency access to single rows from billions of records; enables Random access
Sequential access of data only	Internally uses Hash tables and provides random access; stores the data in indexed HDFS files for faster lookups








When to use Hbase?

when random, real-time read/write access to Big Data is needed

billions of rows x **millions of columns**



Web indexing

Satellite imagery

When to use Hbase?

Need to perform many thousands of operations per second on multiple TB/PB of data

Access patterns are well-know and simple

HBase is

Distributed **column-oriented** database
built on top of the Hadoop file system

Horizontally scalable

**HBase is NOT
a Relational Database!**

HBase is NOT a Relational Database!

HBase is sparse,

*Lots of NULL
empty values*

HBase is NOT a Relational Database!

HBase is sparse,
distributed,

*“Share-nothing”
architecture*

HBase is NOT a Relational Database!

HBase is sparse,
distributed,
persistent,

HBase is NOT a Relational Database!

HBase is sparse,
distributed,
persistent,
multi-dimensional

HBBase is NOT a Relational Database!

HBBase is sparse,
distributed,
persistent,
multi-dimensional
sorted map or **Key/value** store

HBase Data Model

HBBase Data Model

Based on Google's
BigTable

BigTable is a distributed storage system for managing structured data designed to scale to a very large size

Why BigTable?

Challenge - RDBMS performance for very large scale analytic processing

Large scale analytic processing

Big queries – typically range or table scans

Big databases (100s of TB/PB/ZB etc)

Bigtable

Similar to a database

NOT a full relational data model

Data indexed using row and column names

Uses MapReduce

A Bigtable table is:

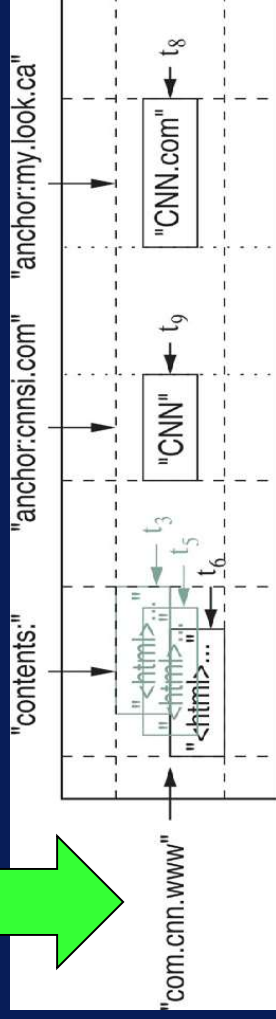

- Sparse
- Distributed
- Persistent
- Multidimensional
- Sorted map

The data in the tables is organized into three dimensions

Rows, Columns, Timestamps

(row:string, column:string, time:int64) → string

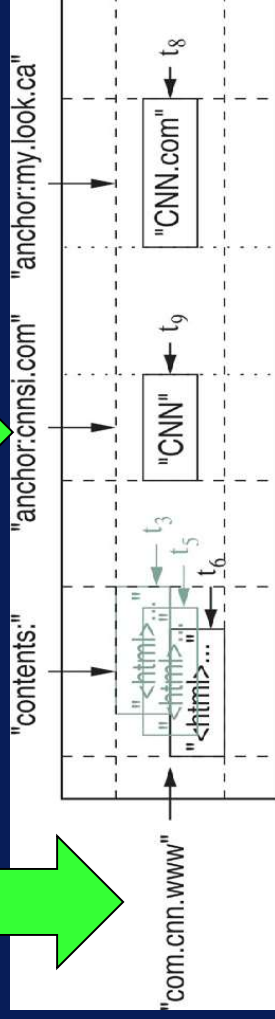


A cell is the storage referenced by a particular
row key,



row:string, column:string, time:int64


A cell is the storage referenced by a particular

row key, column key



row:string, column:string, time:int64

A cell is the storage referenced by a particular
row key, column key and timestamp



From the original Google Bigtable paper

row:string, column:string, time:int64

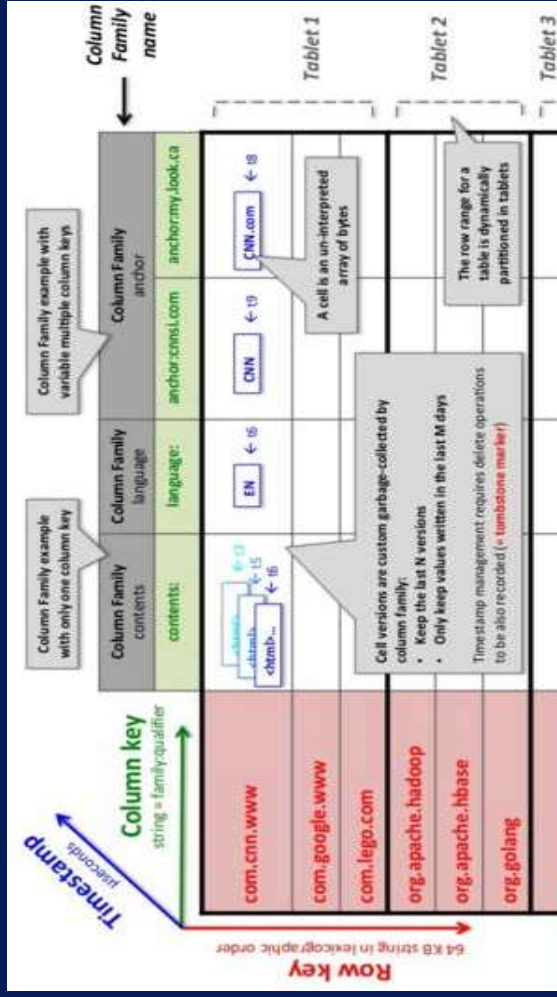
Rows in Big Table

Rows in Bigtable

Table is collection of rows

Bigtable Tablets

Several rows are grouped in tablets
Tablets are distributed
Stored close to each other



Big Table Columns

Column-oriented

The most basic unit in HBase is a column

HBase Columns

Columns could have multiple versions

Distinct values contained in a separate
cells

Columns

One or more columns form a row addressed uniquely by row key

A table is a collection of rows

Column name

A column key is named with

syntax → **family:qualifier**

Columns

Column keys are grouped into column families

Column Families

Semantical boundaries between data

Column families and columns are stored together in the same low-level storage file -> Hfile

HBase: Keys and Column Families

Each row has a Key

Each record is divided into Column Families

row key	personal_data		demographic		...
PersonID	Name	Address	BirthDate	Gender	...
1	H. Houdini	Budapest, Hungary	1926-10-31	M	...
2	D. Copper	New Jersey, USA	1956-09-16	M	...
3	Merlin	Stonehenge, England	1136-12-03	F	...
...
500,000,000	F. Cailliac	Nevada, USA	1964-01-07	M	...

Each column family consists of one or more Columns

Multi-versioned

One column family can have variable no of columns

Cells within a Column family are sorted physically

Very Sparse, most cell has NULL value

Multi-versioned

One column family can have variable no of columns

Cells within a Column family are sorted physically

Very Sparse, most cell has NULL value

Multi-versioned

One column family can have variable no of columns

Cells within a Column family are sorted physically

Very Sparse, most cell has NULL value

Multi-versioned

One column family can have variable no of columns

Cells within a Column family are sorted physically

Very Sparse, most cell has NULL value

Multi-versioned

One column family can have variable no of columns

Cells within a Column family are sorted physically

Very Sparse, most cell has NULL value

Multi-versioned

One column family can have variable no of columns

Cells within a Column family are sorted physically

Very Sparse, most cell has NULL value

New chunk Big Table Timestamps

Timestamps

A cell can hold multiple versions of the data

Timestamps

Timestamps set by Bigtable or
client applications

Timestamps

Data is stored so that new data
are fastest to read

From Bigtable to HBase

- It is open source
- Good integration for the Hadoop
- *No real indexes*
- *Automatic partitioning*


From Bigtable to HBase

- *Scale linearly and automatically with new nodes*
- *Commodity hardware*
- *Fault tolerance*
- *Batch processing*

HBase Data Model


Map indexed by a row key, column key, and a timestamp

(row:string, column:string, time:int64) →
uninterpreted byte array



Data Model

Supports lookups, inserts, deletes
Single row transactions only



Rows and Columns Summary

Rows maintained in sorted lexicographic order

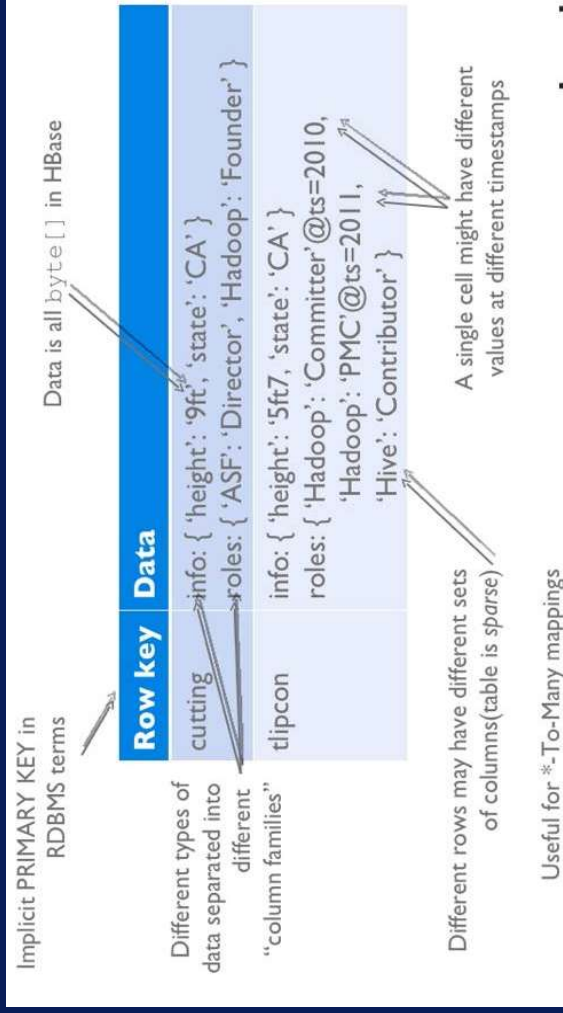
Efficient row scans

Row ranges dynamically partitioned into tablets

Columns grouped into column families

- Column key = *family:qualifier*
- Column families - locality indications
- Boundless number of columns

HBase Example



HBase Data Model

HBase schema consists of several *Tables*

Each table consists of a set of *Column Families*

Columns are not part of the schema

Row key	Data
cutting	info: { 'height': '9ft', 'state': 'CA' } roles: { 'ASF': 'Director', 'Hadoop': 'Founder' }
tlipcon	info: { 'height': '5ft7', 'state': 'CA' } roles: { 'Hadoop': 'Committer'@ts=2010, 'Hadoop': 'PMC'@ts=2011, 'Hive': 'Contributor' }

↑
*"Roles" column family
has different columns in
different cells*

HBase Data Model

HBase has *Dynamic Columns*

Column names are encoded inside the cells

Different cells can have different columns

Row key	Data
cutting	info: { 'height': '9ft', 'state': 'CA' } roles: { 'ASF': 'Director', 'Hadoop': 'Founder' }
tlipcon	info: { 'height': '5ft7', 'state': 'CA' } roles: { 'Hadoop': 'Committer'@ts=2010, 'Hadoop': 'PMC'@ts=2011, 'Hive': 'Contributor' }

“Roles” column family
has different columns in
different cells

CAP Theorem for HBase

HBase provides

Consistency and

Partition Tolerance but is

NOT always **Available**

Accessing HBase

HBase Access

- HBase Shell
- Native Java API
- C/C++ HBase client
- Thrift Server
- REST
- Spark

HBase Shell

Provides interactive commands for manipulating database

- Create/delete tables

- Insert/update/read from tables

- Manage regions

Basic HBase Operations

Get – Retrieves a row of data based on the row key

Put - Inserts a row with data based on the row key

Scan - Finds all matching rows based on the row key

HBase API Operations

get(row) put(row, Map<column, value>)

scan(key range, filter)

increment(row, columns)

Check and Put, delete etc.

Quote all names

Table and column names

Single quotes for text

```
hbase> get 'timestamp1', 'RowId'
```

Double quotes for binary

```
hbase> get 't1', "key\x03\x3f\xcd"
```

Specifying parameters

```
hbase> get 'UserTable', 'userId1', {COLUMN => 'address:str'}
```

HBase Shell Commands

General

Data Definition Language (DDL)

Data Manipulation Language (DML)

Cluster administration

HBase Shell Exercise

Start the QuickStart VM

Open a terminal

At the prompt – type in:

```
[cloudera@quickstart~]$ hbase shell
```

Learn more about each command

- `hbase> help "<command>"`

Display cluster's status via status command

```
hbase> status
```

```
hbase> status 'detailed'
```


Create HBase table

```
hbase> create 'Movies',{NAME=>'info'},{Name=>'director'}
```

```
hbase> put 'Movies','1','info:title','Godfather'
```

```
hbase> put 'Movies','1','info:star','Marlon Brando'
```

```
hbase> put 'Movies','1','info:star','Al Pacino'
```

```
hbase> put 'Movies','1','info:type','Crime'
```

```
hbase> put 'Movies','1','info:type','Drama'
```

Create HBase table

```
hbase> create 'Movies',{NAME=>'info'},{Name=>'director'}

hbase> put 'Movies','1','info:title','Godfather'
hbase> put 'Movies','1','info:star','Marlon Brando'
hbase> put 'Movies','1','info:star','Al Pacino'
hbase> put 'Movies','1','info:movietype','Crime'
hbase> put 'Movies','1','info:movietype','Drama'
hbase> put 'Movies','1','director:First','Francis'
hbase> put 'Movies','1','director:Middle','Ford'
hbase> put 'Movies','1','director:Last','Copola'
```

Create HBase table

```
hbase> create 'Movies',{NAME=>'info'},{Name=>'director'}

hbase> put 'Movies',2,'info:title','Pulp Fiction'
hbase> put 'Movies',2,'info:star','John Travolta'
hbase> put 'Movies',2,'info:star','Uma Thurman'
hbase> put 'Movies',2,'info:star','Samuel Jackson'
hbase> put 'Movies',2,'info:movietype','Drama'
hbase> put 'Movies',2,'director:First','Quentin'
hbase> put 'Movies',2,'director:Last','Tarantino'
```

Verify your data

```
hbase> get 'Movies', '1'
```

```
hbase> get 'Movies', '2'
```

Change data

```
hbase> put 'Movies', '2', 'info:star', 'Samuel  
L. Jackson'
```

```
hbase> get 'Movies', '2'  
hbase> scan 'Movies'
```

Delete data

```
hbase> delete 'Movies','1','info:star'
```

```
hbase>disable 'Movies'
```

```
Hbse>drop 'Movies'
```

HUE Interface Tutorial

HBase Shell Exercise

```
$ hbase shell  
> list
```

*Verify that cluster
is running*

HBase Shell Exercise

```
$ hbase shell
```

```
> create 'test', 'data'
```

*Create a simple
table*

HBase Shell Exercise

```
> put 'test', 'row1', 'data:1', 'value1'
```

*Populate table
with records*

HBase Shell Exercise

```
> put 'test', 'row2', 'data:2', 'value2'
```

HBase Shell Exercise

```
> put 'test', 'row3', 'data:3', 'value3'
```

HBase Shell Exercise

>scan 'test'

*Retrieve rows by
scanning the
entire table*

HBase Shell Exercise

```
> get 'test', 'row2', 'data:2', 'value2'
```

HBase Shell Exercise

```
> delete 'test', 'row2', 'data:2',  
      'value2'
```

HBase Shell Exercise

>scan 'test'

*Retrieve rows by
scanning the
entire table*

HBase Shell Exercise

>disable 'test'

Puts the table "off-line"

*Must disable
before dropping
the table*

HBase Shell Exercise

```
>scan 'test'  
>disable 'test'  
>drop 'test'
```

HBase Shell Exercise

```
>scan 'test'  
>disable 'test'  
>drop 'test'  
>list
```

HBase in Conclusion

HBase vs. RDBMS Revisited

HBase	vs.	RDBMS
Column-oriented		Row oriented
Flexible schema, add columns on the fly		Fixed schema
Good with sparse tables		Not optimized for sparse tables
No query language		SQL
Wide tables		Narrow tables
Joins using MapReduce		Natively performs joins
Tight integration with MapReduce		Minimal if any integration with MapReduce
Horizontal scalability – just add hardware		Hard to shard and scale
De-normalized		Normalized
No transactions		Transactional
Semi-structured & structured data		Structured data

When to consider using HBase?

**Hundreds of millions or
billions of rows**

**Not optimized for classic
transactional applications or
relational analytics**

When to consider using HBase?

If your application has a variable schema where each row is slightly different

Example

Row key	Data
cutting	info: { 'height': '9ft', 'state': 'CA' } roles: { 'ASF': 'Director', 'Hadoop': 'Founder' }
tlipcon	info: { 'height': '5ft7', 'state': 'CA' } roles: { 'Hadoop': 'Committer'@ts=2010, 'Hadoop': 'PMC'@ts=2011, 'Hive': 'Contributor' }

info Column Family

Row key	Column key	Timestamp	Cell value
cutting	info:height	1273516197868	9ft
cutting	info:state	1043871824184	CA
tlipcon	info:height	1273878447049	5ft7
tlipcon	info:state	1273616297446	CA

roles Column Family

Row key	Column key	Timestamp	Cell value
cutting	roles:ASF	1273871823022	Director
cutting	roles:Hadoop	1183746289103	Founder
tlipcon	roles:Hadoop	1300062064923	PMC
tlipcon	roles:Hadoop	1293388212294	Committer
tlipcon	roles:Hive	1273616297446	Contributor

Sorted
on disk by
Row key, Col
key,
descending
timestamp

Milliseconds since unix epoch

cloudera