

# Assignment 1

## Linear Regression

---

**Due Date: 31st March, 2022**

### Guidelines

- Submit all of your code and results in a single zip file with name FirstName\_RollNumber\_04.zip
- Submit single zip file containing a) code (tas01.ipynb, tas) b) report
- There should be a Report.pdf detailing your experience and highlighting any interesting result.
- Email instructor or TAs if there are any questions.

### Overview

In this assignment, you will implement linear regression and get to see it work on data. Linear regression is a linear approach to model the relation b/w a dependent variable and one or more independent variables. Mathematically,

$$\bar{y} = mX + c$$

In the above equation X is an independent variable and Y is dependent, m is slope and c is the y-intercept. The loss is the error in our predicted values of m and c. You will be using Mean Squared error to calculate the loss between true and predicted values.

$$E = \frac{1}{n} \sum_{i=0}^n (y_i - \bar{y}_i)^2$$

Here  $y_i$  is the actual value and  $\bar{y}_i$  is the predicted value and n is the number of training examples. Next step is to find optimal values of m and c using the gradient descent algorithm. Gradient descent algorithm can be broken down to following steps:

- Randomly initialize the values of m and c
- Set a value of learning rate (L), epochs.
- For each epoch:
  - Find  $\bar{y}$  and then compute derivative of E w.r.t to m ( $dE/dm$ ) and c ( $dE/dc$ ).
  - Update the values of m and c using following formulae:

$$m = m - L (dE/dm)$$

$$c = c - L (dE/dc)$$

- Repeat until epoch < epochs.

You will be using a subset of house price prediction dataset available on kaggle. **For now, we would encourage you to implement linear regression from scratch using numpy. Please refrain from using tools like scikit-learn.**

## Task 01: Train/Validation/Test Split (5 Marks)

- Perform train/validation/test (70%/15%/15%) random split on both task01.csv and task02.csv separately. (You are not allowed to use sklearn train\_test\_split function for this task)

## Task 02: Linear Regression with One Variable (25 Marks)

In this task, you will implement linear regression with one variable to predict the apartment price. Download the file **task01.csv**. This file contains the dataset for task 1. The csv file contains two columns Area(sq\_ft) and price and has 1460 entries. Area(sq\_ft) is your independent variable and price is dependent variable.

**Step 1:** Read the data files using pandas.

**Step 2:** Randomly initialize the  $m$  and  $c$ . Set the learning rate to 0.001 and set the number of epochs to 100 (You are encouraged to experiment with a different set of learning rate and number of epochs).

**Step 3:** Write a function *loss\_func(y\_true, y\_pred)* that takes predicted apartment price and original apartment price as input and returns the mean squared error b/w the true and predicted price.

**Step 4:** Write a function *grad\_descent(learning\_rate, epochs, m, c)* that performs the gradient descent to reach optimal values of  $m$  and  $c$  for the given dataset. The algorithm is explained in the overview section. The function should return final values of  $m$ ,  $c$  and an array containing the loss values at each epoch for both training and validation.

**Step 5:** Plot the epoch vs training and validation loss using matplotlib.

**Step 6:** Plot the training data with linear regression fit and find test loss.

## Task 03: Linear Regression with Multiple Variables (60 Marks)

In this task, you will implement linear regression with multiple variables to predict the apartment price. Download the file **task02.csv**. This file contains the dataset for task 2. The csv file contains five columns Area(sq\_ft), Bedrooms, Kitchen, YearBuilt and price and has 1460 entries. Area(sq\_ft), Bedrooms, Kitchen, YearBuilt are your independent variables and price is dependent variable. Zero value in the Bedroom and kitchen column indicates that there is no dedicated area for them (example studio apartment). **For this task you need to report results for both with and without feature normalization.** Normalization is done by subtracting the mean and dividing by standard deviation.

**Step 1:** Read the data file using pandas

**Step 2:** Randomly initialize the  $m$  (Hint: for each feature/variable you will have to compute separate  $m$ ) and  $c$ . Set the learning rate to 0.001 and set the number of epochs to 100 (You are encouraged to experiment with a different set of learning rate and number of epochs).

**Step 3:** Write a function `loss_func(y_true, y_pred)` that takes predicted apartment price and original apartment price as input and returns the mean squared error b/w the true and predicted price.

**Step 4:** Write a function `grad_descent(learning_rate, epochs, m, c)` that performs the gradient descent to reach optimal values of `m` and `c` for the given dataset. The algorithm is explained in the overview section. The function should return final values of `m`, `c` and an array containing the loss values at each epoch.

**Step 5:** Plot the epoch vs training and validation loss values using matplotlib and find test loss.

**For task 2 and 3:** If you are getting NaN values for `m` and `c` on the given parameters try decreasing the learning rate or number of epochs. Another possible option could be to remove outliers from data (you can check z-score, Inter quartile range (IQR) outlier removal method or simply decide outliers by looking at scatter plot). Remember best practice is to remove outliers before train/test/val split.

## Report (10 Marks)

### Task 01:

- Briefly describe the outlier removal technique you used (if any)

### Task 02:

- Describe any important decision you made and any interesting insights you gained.
- Report the parameters which gave the minimum loss along with values of slope, intercept, training loss, validation loss and test loss.
- Comment on training and validation loss curves.

### Task 03:

- Describe any important decision you made and any interesting insights you gained.
- Report the parameters which gave the minimum loss along with values of slope, intercept, training loss, validation loss and test loss.
- Comment on training and validation loss curves.
- Which method gave you the best results with normalization or without normalization? And why? Briefly explain.